

## Supervised Learning: Regression, Part II

Noah Simon & Daniela Witten

July 15-17, 2015  
Summer Institute for Statistics of Big Data  
University of Washington

## Linear Models in High Dimensions

- ▶ When  $p$  is large, least squares regression will lead to very low training error but terrible test error.
- ▶ We will now see some approaches for fitting linear models in high dimensions,  $p \gg n$ .
- ▶ These approaches also work well when  $p \approx n$  or  $n > p$ .

## Motivating example

- ▶ We would like to build a model to predict survival time for breast cancer patients using a number of clinical measurements (tumor stage, tumor grade, tumor size, patient age, etc.) as well as some biomarkers.
- ▶ For instance, these biomarkers could be:
  - ▶ the expression levels of genes measured using a microarray.
  - ▶ protein levels.
  - ▶ mutations in genes potentially implicated in breast cancer.
- ▶ How can we develop a model with low test error in this setting?

## Remember

- ▶ We have  $n$  **training observations**.
- ▶ Our goal is to get a model that will perform well on **future test observations**.
- ▶ We'll incur some bias in order to reduce variance.

## Variable Pre-Selection

The simplest approach for fitting a model in high dimensions:

1. Choose a small set of variables, say the  $q$  variables that are most correlated with the response, where  $q < n$  and  $q < p$ .
2. Use least squares to fit a model predicting  $y$  using only these  $q$  variables.

This approach is simple and straightforward.

## Variable Pre-Selection in R

```
xtr <- matrix(rnorm(100*100),ncol=100)
beta <- c(rep(1,10),rep(0,90))
ytr <- xtr%*%beta + rnorm(100)
cors <- cor(xtr,ytr)
whichers <- which(abs(cors)>.2)
mod <- lm(ytr~xtr[,whichers])
print(summary(mod))
```

## How Many Variable to Use?

- ▶ We need a way to choose  $q$ , the number of variables used in the regression model.
- ▶ We want  $q$  that minimizes the test error.
- ▶ For a range of values of  $q$ , we can perform the validation set approach, leave-one-out cross-validation, or  $K$ -fold cross-validation in order to estimate the test error.
- ▶ Then choose the value of  $q$  for which the estimated test error is smallest.

## Estimating the Test Error For a Given $q$

This is the **right** way to estimate the test error using the validation set approach:

1. Split the observations into a training set and a validation set.
2. Using the training set only:
  - a. Identify the  $q$  variables most associated with the response.
  - b. Use least squares to fit a model predicting  $y$  using those  $q$  variables.
  - c. Let  $\hat{\beta}_1, \dots, \hat{\beta}_q$  denote the resulting coefficient estimates.
3. Use  $\hat{\beta}_1, \dots, \hat{\beta}_q$  obtained on training set to predict response on validation set, and compute the validation set MSE.

## Estimating the Test Error For a Given $q$

This is the **wrong** way to estimate the test error using the validation set approach:

1. Identify the  $q$  variables most associated with the response on the full data set.
2. Split the observations into a training set and a validation set.
3. Using the training set only:
  - a. Use least squares to fit a model predicting  $y$  using those  $q$  variables.
  - b. Let  $\hat{\beta}_1, \dots, \hat{\beta}_q$  denote the resulting coefficient estimates.
4. Use  $\hat{\beta}_1, \dots, \hat{\beta}_q$  obtained on training set to predict response on validation set, and compute the validation set MSE.

9 / 49

## Frequently Asked Questions

- ▶ **Q:** Does it really matter how you estimate the test error?  
**A:** Yes.
- ▶ **Q:** Would anyone make such a silly mistake?  
**A:** Yes.

10 / 49

## A Better Approach

- ▶ The variable pre-selection approach is simple and easy to implement – all you need is a way to calculate correlations, and software to fit a linear model using least squares.
- ▶ But it might not work well: just because a bunch of variables are correlated with the response doesn't mean that when used together in a linear model, they will predict the response well.
- ▶ What we really want to do: pick the  $q$  variables that best predict the response.

11 / 49

## Best Subset Selection

- ▶ We would like to consider all possible models using a subset of the  $p$  predictors.
- ▶ In other words, we'd like to consider all  $2^p$  possible models.
- ▶ This is called **best subset selection**.
- ▶ Unfortunately, this is computationally intractable:
  - ▶ When  $p = 3$ ,  $2^p = 8$ .
  - ▶ When  $p = 6$ ,  $2^p = 64$ .
  - ▶ When  $p = 250$ , there are  $2^{250} \approx 10^{80}$  possible models. According to [www.universetoday.com](http://www.universetoday.com), this is around the number of atoms in the known universe.
  - ▶ Not feasible to consider so many models!
- ▶ Need an efficient way to sift through all of these models:  
**forward stepwise regression.**

12 / 49


## Forward Stepwise Regression

1. Use least squares to fit  $p$  univariate regression models, and select the predictor corresponding to the best model (according to e.g. training set MSE).
2. Use least squares to fit  $p - 1$  models containing that one predictor, and each of the  $p - 1$  other predictors. Select the predictors in the best two-variable model.
3. Now use least squares to fit  $p - 2$  models containing those two predictors, and each of the  $p - 2$  other predictors. Select the predictors in the best three-variable model.
4. And so on....

This gives us a nested set of models, containing the predictors

$$\mathcal{M}_1 \subseteq \mathcal{M}_2 \subseteq \mathcal{M}_3 \subseteq \dots$$

## Forward Stepwise Regression With $p = 10$



## Example in R

```
xtr <- matrix(rnorm(100*100),ncol=100)
beta <- c(rep(1,10),rep(0,90))
ytr <- xtr%*%beta + rnorm(100)
library(leaps)
out <- regsubsets(xtr,ytr,nvmax=30,method="forward")
print(summary(out))
print(coef(out,1:10))
```

## Which Value of $q$ is Best?

- This procedure traces out a set of models, containing between 1 and  $p$  variables.
- The  $q$ th model contains  $q$  variables, given by the set  $\mathcal{M}_q$ .
- **Q:** Which value of  $q$  is best?  
**A:** *The one that minimizes the test error!*
- We can select the value of  $q$  using cross-validation or the validation set approach.

## Drawback of Forward Stepwise Selection

- ▶ Forward stepwise selection isn't guaranteed to give you the **best** model containing  $q$  variables.
- ▶ To get the **best** model with  $q$  variables, you'd need to consider every possible one; computationally intractable.
- ▶ For instance, suppose that the best model with one variable is

$$y = \beta_3 X_3 + \epsilon$$

and the best model with two variables is

$$y = \beta_4 X_4 + \beta_8 X_8 + \epsilon.$$

Then forward stepwise selection will not identify the best two-variable model.

- ▶ **Q:** Does this really happen in practice?  
**A:** Yes.

17 / 49

## How To Do Forward Stepwise?

**Wrong:** Split the data into a training set and a validation set. Perform forward stepwise on the training set, and identify the model with best performance on the validation set. Then, refit the model (using those  $q$  variables) on the full data set.

**Right:** Split the data into a training set and a validation set. Perform forward stepwise on the training set, and identify the value of  $q$  corresponding to the best-performing model on the validation set. Then, perform forward stepwise selection in order to obtain a  $q$ -variable model on the full data set.

**Bottom Line:** We estimate the test error in order to choose the correct level of **model complexity**. Then we refit the model on the full data set.

18 / 49

## Let's Try It Out in R!

Chapter 6 R Lab, Part 1  
[www.statlearning.com](http://www.statlearning.com)

19 / 49

## Ridge Regression and the Lasso

- ▶ Forward stepwise selection does a discrete search through model space, considering subsets of the predictors, and fitting each of the resulting models using least squares. Model complexity is controlled by using subsets of the predictors.
- ▶ **Ridge regression** and the **lasso** instead control model complexity by using an alternative to least squares, by **shrinking the regression coefficients**.
- ▶ This is known as **regularization** or **penalization**.
- ▶ **Hot area** in statistical machine learning today.

20 / 49

## Crazy Coefficients

- ▶ When  $p > n$ , some of the variables are **highly correlated**.
- ▶ Why does correlation matter?
  - ▶ Suppose that  $X_1$  and  $X_2$  are highly correlated with each other... assume  $X_1 = X_2$  for the sake of argument.
  - ▶ And suppose that the least squares model is

$$\hat{y} = X_1 - 2X_2 + 3X_3.$$

- ▶ Then this is **also** a least squares model:

$$\hat{y} = 100000001X_1 - 100000002X_2 + 3X_3.$$

- ▶ **Bottom Line:** When there are too many variables, the least squares coefficients can get crazy!
- ▶ This craziness is **directly responsible for poor test error**.
- ▶ It amounts to **too much model complexity**.

21 / 49

## A Solution: Don't Let the Coefficients Get Too Crazy

- ▶ Recall that least squares involves finding  $\beta$  that minimizes

$$\|\mathbf{y} - \mathbf{X}\beta\|^2.$$

- ▶ Ridge regression involves finding  $\beta$  that minimizes

$$\|\mathbf{y} - \mathbf{X}\beta\|^2 + \lambda \sum_j \beta_j^2.$$

- ▶ Equivalently, find  $\beta$  that minimizes

$$\|\mathbf{y} - \mathbf{X}\beta\|^2$$

subject to the constraint that

$$\sum_{j=1}^p \beta_j^2 \leq s.$$

22 / 49

## Ridge Regression


- ▶ Ridge regression coefficient estimates minimize

$$\|\mathbf{y} - \mathbf{X}\beta\|^2 + \lambda \sum_j \beta_j^2.$$

- ▶ Here  $\lambda$  is a nonnegative **tuning parameter** that shrinks the coefficient estimates.
- ▶ When  $\lambda = 0$ , then ridge regression is just the same as least squares.
- ▶ As  $\lambda$  increases, then  $\sum_{j=1}^p (\hat{\beta}_{\lambda,j}^R)^2$  decreases – i.e. coefficients become shrunken towards zero.
- ▶ When  $\lambda = \infty$ ,  $\hat{\beta}_{\lambda}^R = 0$ .

23 / 49

## Ridge Regression As $\lambda$ Varies



24 / 49

## Ridge Regression In Practice

- ▶ Perform ridge regression for a very fine grid of  $\lambda$  values.
- ▶ Use cross-validation or the validation set approach to select the optimal value of  $\lambda$  – that is, the best level of model complexity.
- ▶ Perform ridge on the full data set, using that value of  $\lambda$ .


25 / 49

## Example in R

```
xtr <- matrix(rnorm(100*100),ncol=100)
beta <- c(rep(1,10),rep(0,90))
ytr <- xtr%*%beta + rnorm(100)
library(glmnet)
cv.out <- cv.glmnet(xtr,ytr,alpha=0,nfolds=5)
print(cv.out$cvm)
plot(cv.out)
cat("CV Errors", cv.out$cvm,fill=TRUE)
cat("Lambda with smallest CV Error",
    cv.out$lambda[which.min(cv.out$cvm)],fill=TRUE)
cat("Coefficients", as.numeric(coef(cv.out)),fill=TRUE)
cat("Number of Zero Coefficients",
    sum(abs(coef(cv.out))<1e-8),fill=TRUE)
```

26 / 49

## R Output



27 / 49

## Drawbacks of Ridge

- ▶ Ridge regression is a simple idea and has a number of attractive properties: for instance, you can continuously control model complexity through the tuning parameter  $\lambda$ .
- ▶ But it suffers in terms of model interpretability, since the final model contains **all  $p$  variables, no matter what**.
- ▶ Often want a simpler model involving a subset of the features.
- ▶ **The lasso** involves performing a little tweak to ridge regression so that the resulting model contains **mostly zeros**.
- ▶ In other words, the resulting model is **sparse**. We say that the lasso performs **feature selection**.
- ▶ The lasso is a very active area of research interest in the statistical community!

28 / 49

## The Lasso

- ▶ The lasso involves finding  $\beta$  that minimizes

$$\|y - X\beta\|^2 + \lambda \sum_j |\beta_j|.$$

- ▶ Equivalently, find  $\beta$  that minimizes

$$\|y - X\beta\|^2$$

subject to the constraint that

$$\sum_{j=1}^p |\beta_j| \leq s.$$

- ▶ So lasso is just like ridge, except that  $\beta_j^2$  has been replaced with  $|\beta_j|$ .


29 / 49

## The Lasso

- ▶ Lasso is a lot like ridge:
  - ▶  $\lambda$  is a nonnegative tuning parameter that controls model complexity.
  - ▶ When  $\lambda = 0$ , we get least squares.
  - ▶ When  $\lambda$  is very large, we get  $\hat{\beta}_\lambda^L = 0$ .
- ▶ But unlike ridge, **lasso will give some coefficients exactly equal to zero for intermediate values of  $\lambda$ !**

30 / 49

## Lasso As $\lambda$ Varies



31 / 49

## Lasso In Practice

- ▶ Perform lasso for a very fine grid of  $\lambda$  values.
- ▶ Use cross-validation or the validation set approach to select the optimal value of  $\lambda$  – that is, the best level of model complexity.
- ▶ Perform the lasso on the full data set, using that value of  $\lambda$ .

32 / 49




## Example in R

```
xtr <- matrix(rnorm(100*100),ncol=100)
beta <- c(rep(1,10),rep(0,90))
ytr <- xtr%*%beta + rnorm(100)
library(glmnet)
cv.out <- cv.glmnet(xtr,ytr,alpha=1,nfolds=5)
print(cv.out$cvm)
plot(cv.out)
cat("CV Errors", cv.out$cvm,fill=TRUE)
cat("Lambda with smallest CV Error",
    cv.out$lambda[which.min(cv.out$cvm)],fill=TRUE)
cat("Coefficients", as.numeric(coef(cv.out)),fill=TRUE)
cat("Number of Zero Coefficients",sum(abs(coef(cv.out))<1e-8),
    fill=TRUE)
```


33 / 49

## R Output



34 / 49

## Ridge and Lasso: A Geometric Interpretation



35 / 49

## Let's Try It Out in R!

Chapter 6 R Lab, Part 2  
[www.statlearning.com](http://www.statlearning.com)

36 / 49

## Review

- ▶ So far we have seen two approaches that select subsets of the features and fit a least squares model:
  - ▶ Variable Pre-Selection
  - ▶ Forward Stepwise Selection
- ▶ And we have seen two approaches that fit a shrunk model instead of using least squares:
  - ▶ Ridge regression
  - ▶ Lasso
- ▶ Now we see one final approach, **principal components regression**, that first finds a **low-dimensional subspace** of the data and then fits a model on that low-dimensional subspace, using least squares.

37 / 49

## Principal Components Regression

- ▶ Our data consist of  $n$  observations in a  $p$ -dimensional space.
- ▶ However, not all of those  $p$  dimensions are equally useful, especially when  $p \gg n$ .
- ▶ Many are either completely redundant (correlated features) or uninformative (noise features).
- ▶ Can we find a low-dimensional representation of the variables that captures most of the variability in the data?
- ▶ This is a **dimension reduction** approach.

38 / 49

## PCR

- ▶ Let  $Z_1, Z_2, \dots, Z_M$  represent  $M < p$  **linear combinations** of the  $p$  predictors:

$$Z_m = \sum_{j=1}^p \phi_{mj} X_j.$$


- ▶ Use least squares to fit the model

$$y_i = \theta_0 + \sum_{m=1}^M \theta_m Z_{im} + \epsilon_i, \quad i = 1, \dots, n.$$

- ▶ In other words, we perform least squares using  $M$  new predictors,  $Z_1, \dots, Z_M$ .
- ▶  $Z_1, \dots, Z_M$  chosen to be the **principal components** of the data.

39 / 49

## Principal Components, Conceptually



- ▶ PCs are the linear combinations of the variables that **contain as much as possible of the variability in the features**.
- ▶ Will be discussed further in SISBID Module 4 — Unsupervised Learning.

40 / 49

## PCR

Our final model is linear in the original predictors:

$$\begin{aligned} y_i &= \theta_0 + \sum_{m=1}^M \theta_m Z_{im} + \epsilon_i \\ &= \theta_0 + \sum_{m=1}^M \theta_m \sum_{j=1}^p \phi_{mj} X_{ij} + \epsilon_i \\ &= \theta_0 + \sum_{j=1}^p \left( \sum_{m=1}^M \theta_m \phi_{mj} \right) X_{ij} + \epsilon_i \end{aligned}$$

## More on PCR

- ▶ PCR doesn't yield feature selection – all of the original predictors are involved in the final model.
- ▶ But when  $M$  is small, then PCR can avoid overfitting and can give good results.
- ▶ Choose  $M$  by cross-validation or validation set approach.
- ▶ With  $M = p$ , we just get least squares regression: no dimension reduction occurs!
- ▶ Turns out that PCR is closely related to ridge regression.
- ▶ Shortcoming of PCR: the first  $M$  principal components are guaranteed to explain a lot of the variation in the features, but that doesn't mean that they are predictive of the response!
- ▶ In SISBID Module 4, will see how principal components can be used for **unsupervised learning**.

## Example in R

```
xtr <- matrix(rnorm(100*100),ncol=100)
beta <- c(rep(1,10),rep(0,90))
ytr <- xtr%*%beta + rnorm(100)
library(pls)
out <- pcr(ytr~xtr,scale=TRUE,validation="CV")
summary(out)
validationplot(out,val.type="MSEP")
```

## Let's Try It Out in R!

Chapter 6 R Lab, Part 3  
[www.statlearning.com](http://www.statlearning.com)

## Pros/Cons of Each Approach

Approach	Simplicity?*	Sparsity?**	Predictions?***
Pre-Selection	Good	Yes	So-So
Forward Stepwise	Good	Yes	So-So
Ridge	Medium	No	Great
Lasso	Bad	Yes	Great
PCR	Medium	No	Great

\* How simple is this model-fitting procedure? If you were stranded on a desert island with pretty limited statistical software, could you fit this model?

\*\* Does this approach perform feature selection, i.e. is the resulting model sparse?

\*\*\* How good are the predictions resulting from this model?

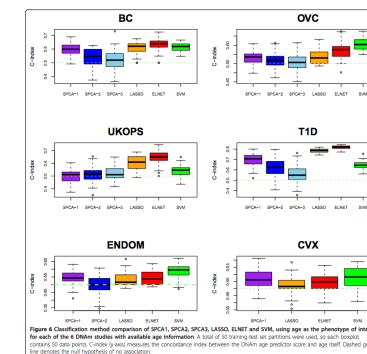
## No “Best” Approach

- ▶ There is no “best” approach to regression in high dimensions.
- ▶ Some approaches will work better than others. For instance:
  - ▶ Lasso will work well if it's really true that just a few features are associated with the response.
  - ▶ Ridge will do better if all of the features are associated with the response.
- ▶ If somebody tells you that one approach is “best”... then they are mistaken. Politely contradict them.
- ▶ While no approach is “best”, some approaches are wrong (e.g.: there is a wrong way to do cross-validation)!

## Predicting Age Using DNA Methylation Data

- ▶ Comparison on 6 data sets
- ▶ SPC: Like principal components regression, but using a subset of features most associated with response. Between 1 and 3 principal components were used.
- ▶ Elastic Net: A hybrid between ridge and lasso.
- ▶ SVM: We'll see it next lecture in the classification context.
- ▶ Citation: Zhuang et al., BMC Bioinformatics, 2012

## Didn't I Tell You? No Best Method!



High C-index indicates a low test error.

## Bottom Line

Much more important than what model you fit is how you fit it.

- ▶ Was cross-validation performed properly?
- ▶ Did you select a model (or level of model complexity) based on an estimate of test error?