

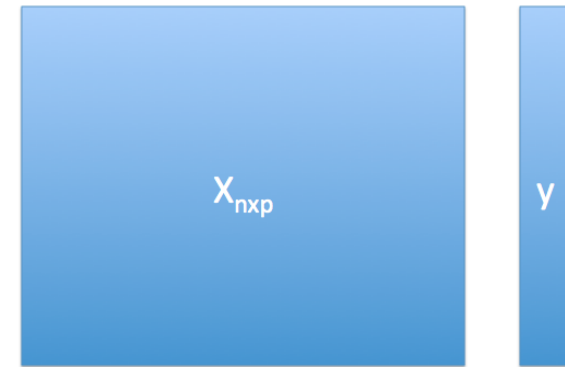
## Supervised Learning: Regression, Part I

Noah Simon & Daniela Witten

July 15-17, 2015  
Summer Institute for Statistics of Big Data  
University of Washington

1 / 43

## Supervised Learning



2 / 43

## Regression Versus Classification

- ▶ **Regression:** Predict a **quantitative** response, such as
  - ▶ blood pressure
  - ▶ cholesterol level
  - ▶ tumor size
- ▶ **Classification:** Predict a **categorical** response, such as
  - ▶ tumor versus normal tissue
  - ▶ heart disease versus no heart disease
  - ▶ subtype of glioblastoma
- ▶ This lecture: **Regression**.

3 / 43

## Linear Models


- ▶ We have  $n$  observations, for each of which we have  $p$  predictor measurements and a response measurement.
- ▶ Want to develop a model of the form

$$y_i = \beta_0 + \beta_1 X_{i1} + \dots + \beta_p X_{ip} + \epsilon_i.$$

- ▶ Here  $\epsilon_i$  is a noise term associated with the  $i$ th observation.
- ▶ Must estimate  $\beta_0, \beta_1, \dots, \beta_p$  — i.e. we must **fit the model**.

4 / 43

## Linear Model With $p = 2$ Predictors



5 / 43

## What Makes a Model Linear?

- ▶ A linear model is **linear in the regression coefficients**!
- ▶ This is a linear model:

$$y_i = \beta_1 \sin(X_{i1}) + \beta_2 X_{i2} X_{i3} + \epsilon_i.$$

- ▶ This is not a linear model:

$$y_i = \beta_1^{X_{i1}} + \sin(\beta_2 X_{i2}) + \epsilon_i.$$

6 / 43

## Linear Models in Matrix Form

- ▶ For simplicity, ignore the intercept  $\beta_0$ .
  - ▶ Assume  $\sum_{i=1}^n y_i = \sum_{i=1}^n X_{ij} = 0$ ; in this case,  $\beta_0 = 0$ .
  - ▶ Alternatively, let the first column of  $\mathbf{X}$  be a column of 1's.
- ▶ In matrix form, we can write the linear model as

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon},$$

i.e.

$$\begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} = \begin{pmatrix} X_{11} & X_{12} & \dots & X_{1p} \\ X_{21} & X_{22} & \dots & X_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ X_{n1} & X_{n2} & \dots & X_{np} \end{pmatrix} \begin{pmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{pmatrix} + \begin{pmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{pmatrix}.$$

7 / 43

## Least Squares Regression

- ▶ There are a lot of ways we could fit the model

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}.$$

- ▶ Most common approach in classical statistics is **least squares**:

$$\underset{\boldsymbol{\beta}}{\text{minimize}} \{ \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|^2 \}.$$

Here  $\|\mathbf{a}\|^2 \equiv \sum_{i=1}^n a_i^2$ .

- ▶ We are looking for  $\beta_1, \dots, \beta_p$  such that

$$\sum_{i=1}^n (y_i - (\beta_1 X_{i1} + \dots + \beta_p X_{ip}))^2$$

is as small as possible.


- ▶ Equivalently, we're looking for coefficient estimates such that

$$\sum_{i=1}^n (y_i - \hat{y}_i)^2$$

is as small as possible, where  $\hat{y}_i$  is the  $i$ th predicted value.

8 / 43

## Least Squares



- Horizontal axis: predictor
  - Vertical axis: response
  - Red dots: observations
  - Purple line: least squares line
- Purple line minimizes sum of squared lengths of the gray lines.

9 / 43

## Let's Try Out Least Squares in R!

Chapter 3 R lab  
[www.statlearning.com](http://www.statlearning.com)

10 / 43

## Least Squares Regression

- When we fit a model, we use a **training set** of observations.
- We get coefficient estimates  $\hat{\beta}_1, \dots, \hat{\beta}_p$ .
- We also get predictions using our model, of the form

$$\hat{y}_i = \hat{\beta}_1 X_{i1} + \dots + \hat{\beta}_p X_{ip}.$$

- We can evaluate the **training error**, i.e. the extent to which the model fits the observations used to train it.
- One way to quantify the training error is using the **mean squared error (MSE)**:

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{1}{n} \sum_{i=1}^n (y_i - (\hat{\beta}_1 X_{i1} + \dots + \hat{\beta}_p X_{ip}))^2.$$

- The training error is closely related to the  $R^2$  for a linear model – that is, the **proportion of variance explained**.
- Big  $R^2 \Leftrightarrow$  Small Training Error.

11 / 43


## Least Squares as More Variables are Included in the Model

- Training error and  $R^2$  are not good ways to evaluate a model's performance, because **they will always improve as more variables are added into the model**.
- The problem? Training error and  $R^2$  evaluate the model's performance on the **training observations**.
- If I had an unlimited number of features to use in developing a model, then I could surely come up with a regression model that **fits the training data perfectly!** Unfortunately, this model wouldn't capture the true signal in the data.
- We really care about the model's performance on **test observations** – observations not used to fit the model.

12 / 43

## The Problem

As we add more variables into the model...



... the training error decreases and the  $R^2$  increases!

13 / 43

## Why is this a Problem?

- ▶ We really care about the model's performance on **observations not used to fit the model!**
  - ▶ Want to predict the survival time of a new patient who walks into the clinic!
  - ▶ Want to diagnose cancer for a patient not used in model training!
  - ▶ Want to predict risk of diabetes for a patient who wasn't used to fit the model!
- ▶ What we really care about:

$$(y_{test} - \hat{y}_{test})^2,$$

where

$$\hat{y}_{test} = \hat{\beta}_1 X_{test,1} + \dots + \hat{\beta}_p X_{test,p},$$


and  $(X_{test}, y_{test})$  **was not used to train the model.**

- ▶ The **test error** is the average of  $(y_{test} - \hat{y}_{test})^2$  over a bunch of test observations.

14 / 43

## Training Error versus Test Error

As we add more variables into the model...




... the training error decreases and the  $R^2$  increases!

**But the test error might not!**

15 / 43

## Training Error versus Test Error

As we add more variables into the model...



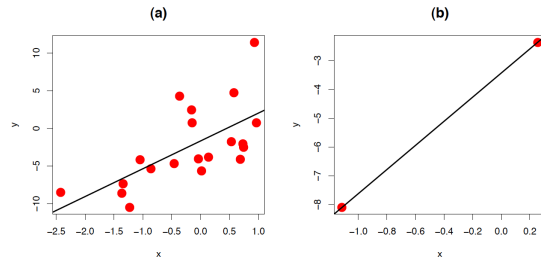
... the training error decreases and the  $R^2$  increases!

**But the test error might not!**

16 / 43

## Why the Number of Variables Matters

- ▶ Linear regression will have a very low training error if  $p$  is large relative to  $n$ .
- ▶ A simple example:



- ▶ When  $n \leq p$ , you can always get a perfect model fit to the training data!
- ▶ But the test error will be awful.

17 / 43

## Model Complexity, Training Error, and Test Error

- ▶ In this course, we will consider various types of models.
- ▶ We will be very concerned with **model complexity**: e.g. the number of variables used to fit a model.
- ▶ As we fit more complex models – e.g. models with more variables – the training error will always decrease.
- ▶ But the test error might not.
- ▶ As we will see, the number of variables in the model is not the only – or even the best – way to quantify model complexity.


18 / 43

## An Example In R

```
xtr <- matrix(rnorm(100*100),ncol=100)
xte <- matrix(rnorm(100000*100),ncol=100)
beta <- c(rep(1,10),rep(0,90))
ytr <- xtr*beta + rnorm(100)
yte <- xte*beta + rnorm(100000)
rsq <- trainerr <- testerr <- NULL
for(i in 2:100){
  mod <- lm(ytr~xtr[,1:i])
  rsq <- c(rsq,summary(mod)$r.squared)
  beta <- mod$coef[-1]
  intercept <- mod$coef[1]
  trainerr <- c(trainerr, mean((xtr[,1:i]*beta+intercept - ytr)^2))
  testerr <- c(testerr, mean((xte[,1:i]*beta+intercept - yte)^2))
}
par(mfrow=c(1,3))
plot(2:100,rsq, xlab='Number of Variables', ylab="R Squared", log="y")
abline(v=10,col="red")
plot(2:100,trainerr, xlab='Number of Variables', ylab="Training Error",log="y")
abline(v=10,col="red")
plot(2:100,testerr, xlab='Number of Variables', ylab="Test Error",log="y")
abline(v=10,col="red")
```

19 / 43

## Output of R Code



- ▶ 1st 10 variables are related to response; remaining 90 are not.
- ▶  $R^2$  increases and training error decreases as more variables are added to the model.
- ▶ Test error is lowest when only signal variables in model.

20 / 43

## Bias and Variance


- ▶ As model complexity increases, the **bias** of  $\hat{\beta}$  – the average difference between  $\beta$  and  $\hat{\beta}$ , if we were to repeat the experiment a huge number of times – will decrease.
- ▶ But as complexity increases, the **variance** of  $\hat{\beta}$  – the amount by which the  $\hat{\beta}$ 's will differ across experiments – will increase.
- ▶ The test error depends on both the bias and variance:

$$\text{Test Error} = \text{Bias}^2 + \text{Variance}.$$

- ▶ There is a **bias-variance trade-off**. We want a model that is sufficiently complex as to have not too much bias, but not so complex that it has too much variance.

21 / 43

## A Really Fundamental Picture



22 / 43

## Overfitting

- ▶ Fitting an overly complex model – a model that has too much variance – is known as **overfitting**.
- ▶ When  $p \approx n$  or  $p \gg n$ , must work hard to avoid overfitting.
- ▶ In particular, we must rely not on training error, but on test error, as a measure of model performance.
- ▶ How can we estimate the test error?
  1. The validation set approach.
  2. Leave-one-out cross-validation.
  3.  $K$ -fold cross-validation.

23 / 43

## Three Ways to Estimate Test Error

1. **The validation set approach.**
2. Leave-one-out cross-validation.
3.  $K$ -fold cross-validation.

24 / 43

## Validation Set Approach

Split the  $n$  observations into two sets of approximately equal size. Train on one set, and evaluate performance on the other.



25 / 43

## Validation Set Approach

1. Split the observations into two sets of approximately equal size, a **training set** and a **validation set**.
  - a. Fit the model using the training observations. Let  $\hat{\beta}_{(train)}$  denote the regression coefficient estimates.
  - b. For each observation in the validation set, compute  $e_i = (y_i - \mathbf{x}_i^T \hat{\beta}_{(train)})^2$ .
2. Calculate the total validation set error by summing the  $e_i$ 's over all of the validation set observations.

26 / 43


## Three Ways to Estimate Test Error

1. The validation set approach.
2. **Leave-one-out cross-validation**.
3.  $K$ -fold cross-validation.

27 / 43

## Leave-One-Out Cross-Validation

Fit  $n$  models, each on  $n - 1$  of the observations. Evaluate each model on the left-out observation.



28 / 43

## Leave-One-Out Cross-Validation

- For  $i = 1, \dots, n$ :
  - Fit the model using observations  $1, \dots, i-1, i+1, \dots, n$ . Let  $\hat{\beta}_{(i)}$  denote the regression coefficient estimates.
  - Compute  $e_i = (y_i - \mathbf{x}_i^T \hat{\beta}_{(i)})^2$ .
- Calculate  $\sum_{i=1}^n e_i$ , the total CV error.

29 / 43


## Three Ways to Estimate Test Error

- The validation set approach.
- Leave-one-out cross-validation.
- K-fold cross-validation.**

30 / 43

## 5-Fold Cross-Validation

Split the observations into 5 sets. Repeatedly train the model on 4 sets and evaluate its performance on the 5th.



31 / 43

## K-fold cross-validation

A generalization of leave-one-out cross-validation:

- Split the  $n$  observations into  $K$  equally-sized folds.
- For  $k = 1, \dots, K$ :
  - Fit the model using the observations **not** in the  $k$ th fold.
  - Let  $e_k$  denote the test error for the observations in the  $k$ th fold.
- Calculate  $\sum_{k=1}^K e_k$ , the total CV error.

32 / 43




## An Example In R

```
xtr <- matrix(rnorm(100*100),ncol=100)
beta <- c(rep(1,10),rep(0,90))
ytr <- xtr%*%beta + rnorm(100)
cv.err <- NULL
for(i in 2:50){
  dat <- data.frame(x=xtr[,1:i],y=ytr)
  mod <- glm(y~.,data=dat)
  cv.err <- c(cv.err, cv.glm(dat,mod,K=6)$delta[1])
}
plot(2:50, cv.err, xlab="Number of Variables",
     ylab="6-Fold CV Error", log="y")
abline(v=10, col="red")
```

33 / 43

## Output of R Code



- Six-fold CV identifies the model with just over ten predictors.
- First ten predictors contain signal, and the rest are noise.

34 / 43

## After Estimating the Test Error...

After we estimate the test error, we refit the “best” model on all of the available observations.

35 / 43

## Let's Try Out Cross-Validation in R!

Chapter 5 R lab  
First Half: Cross-Validation  
[www.statlearning.com](http://www.statlearning.com)

36 / 43

## Regression with Big Data

- ▶ We usually cannot perform least squares regression to fit a model with high-dimensional data, because we will get zero training error but a terrible test error.
- ▶ Instead, we must fit a **less complex model**, e.g. a model with **fewer variables**.
- ▶ We will consider five ways to fit less complex models:
  1. Variable Pre-Selection
  2. Forward Stepwise Regression
  3. Ridge Regression
  4. Lasso Regression
  5. Principal Components Regression
- ▶ These are **alternatives to least squares**.
- ▶ Each of these approaches will allow us to choose the **level of complexity** – e.g. the number of variables in the model.
- ▶ Will select level of complexity using cross-validation or validation set approach.

37 / 43

## The Fundamental Truth About High-Dimensional Data

If you


- ▶ fit your model carelessly;
- ▶ do not properly estimate the test error;
- ▶ or select a model based on training error;

then you **will woefully overfit your training data**, leading to a model that looks good on training data but will perform atrociously on future observations.

Our intuition **breaks down** in high dimensions, and so rigorous model-fitting is crucial.

38 / 43

## The Curse of Dimensionality



39 / 43

## The Curse of Dimensionality

**Q:** A data set with more variables is better than a data set with fewer variables, right?

**A:** Not necessarily!

**Noise variables** – such as genes whose expression levels are not truly associated with the response being studied – will simply increase the risk of overfitting, and the difficulty of developing an effective model that will perform well on future observations.

On the other hand, more **signal variables** – variables that are truly associated with the response being studied – are always useful!

40 / 43

## Every Biostatisticians' Favorite Anecdote

A biostatistician walks into a collaborator's office with a list of genes found to be predictive of survival time in a condition of interest....

41 / 43

## Wise Words

Common mistakes are simple, and simple mistakes are common.

– Keith Baggerly (Instructor for SISBID Module 5)

42 / 43

## Before You're Done Your Analysis

- ▶ Estimate the test error.
- ▶ Do a “sanity check” whenever possible.
  - ▶ “Spot-check” the variables that have the largest coefficients in the model.
  - ▶ Rewrite your code from scratch. Do you get the same answer again?

43 / 43